

# Lab 3: Decision Support System for Cardiovascular disease

```
In [1]: from google.colab import drive  
import pandas as pd  
import seaborn as sns  
import numpy as np  
import matplotlib.pyplot as plt
```

```
In [2]: drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force\_remount=True).

```
In [3]: # Use path  
dataset = pd.read_csv('/content/drive/MyDrive/heart_disease.csv')
```

```
In [4]: dataset.head(5)
```

Out[4]:

|   | HeartDisease | BMI   | Smoking | AlcoholDrinking | Stroke | PhysicalHealth | MentalHealth | I    |
|---|--------------|-------|---------|-----------------|--------|----------------|--------------|------|
| 0 | No           | 16.60 | Yes     |                 | No     | No             | 3.0          | 30.0 |
| 1 | No           | 20.34 | No      |                 | No     | Yes            | 0.0          | 0.0  |
| 2 | No           | 26.58 | Yes     |                 | No     | No             | 20.0         | 30.0 |
| 3 | No           | 24.21 | No      |                 | No     | No             | 0.0          | 0.0  |
| 4 | No           | 23.71 | No      |                 | No     | No             | 28.0         | 0.0  |

```
In [5]: dataset.isna().sum()
```

Out[5]:

|                         | 0   |
|-------------------------|-----|
| <b>HeartDisease</b>     | 0   |
| <b>BMI</b>              | 6   |
| <b>Smoking</b>          | 74  |
| <b>AlcoholDrinking</b>  | 115 |
| <b>Stroke</b>           | 21  |
| <b>PhysicalHealth</b>   | 17  |
| <b>MentalHealth</b>     | 14  |
| <b>DiffWalking</b>      | 33  |
| <b>Sex</b>              | 313 |
| <b>AgeCategory</b>      | 44  |
| <b>Race</b>             | 39  |
| <b>Diabetic</b>         | 33  |
| <b>PhysicalActivity</b> | 16  |
| <b>GenHealth</b>        | 493 |
| <b>SleepTime</b>        | 219 |
| <b>Asthma</b>           | 0   |
| <b>KidneyDisease</b>    | 17  |
| <b>SkinCancer</b>       | 0   |

**dtype:** int64

In [6]:

```
dataset['BMI'].fillna(dataset['BMI'].mean(), inplace=True)
dataset['Smoking'].fillna(dataset['Smoking'].mode()[0], inplace=True)
dataset['AlcoholDrinking'].fillna(dataset['AlcoholDrinking'].mode()[0], inplace=True)
dataset['Stroke'].fillna(dataset['Stroke'].mode()[0], inplace=True)
dataset['PhysicalHealth'].fillna(dataset['PhysicalHealth'].mean(), inplace=True)
dataset['MentalHealth'].fillna(dataset['MentalHealth'].mean(), inplace=True)
dataset['DiffWalking'].fillna(dataset['DiffWalking'].mode()[0], inplace=True)
dataset['Sex'].fillna(dataset['Sex'].mode()[0], inplace=True)
dataset['AgeCategory'].fillna(dataset['AgeCategory'].mode()[0], inplace=True)
dataset['Race'].fillna(dataset['Race'].mode()[0], inplace=True)
dataset['Diabetic'].fillna(dataset['Diabetic'].mode()[0], inplace=True)
dataset['PhysicalActivity'].fillna(dataset['PhysicalActivity'].mode()[0], inplace=True)
dataset['GenHealth'].fillna(dataset['GenHealth'].mode()[0], inplace=True)
dataset['SleepTime'].fillna(dataset['SleepTime'].mean(), inplace=True)
dataset['KidneyDisease'].fillna(dataset['KidneyDisease'].mode()[0], inplace=True)
```

```
/tmp/ipython-input-2612038968.py:1: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method. The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.
```

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
dataset['BMI'].fillna(dataset['BMI'].mean(), inplace=True)  
/tmp/ipython-input-2612038968.py:2: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method. The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.
```

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
dataset['Smoking'].fillna(dataset['Smoking'].mode()[0], inplace=True)  
/tmp/ipython-input-2612038968.py:3: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method. The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.
```

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
dataset['AlcoholDrinking'].fillna(dataset['AlcoholDrinking'].mode()[0], inplace=True)  
/tmp/ipython-input-2612038968.py:4: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method. The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.
```

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
dataset['Stroke'].fillna(dataset['Stroke'].mode()[0], inplace=True)  
/tmp/ipython-input-2612038968.py:5: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method. The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.
```

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
dataset['PhysicalHealth'].fillna(dataset['PhysicalHealth'].mean(), inplace=True)
```

```
/tmp/ipython-input-2612038968.py:6: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method. The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.
```

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
dataset['MentalHealth'].fillna(dataset['MentalHealth'].mean(), inplace=True)  
/tmp/ipython-input-2612038968.py:7: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method. The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.
```

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
dataset['DiffWalking'].fillna(dataset['DiffWalking'].mode()[0], inplace=True)  
/tmp/ipython-input-2612038968.py:8: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method. The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.
```

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
dataset['Sex'].fillna(dataset['Sex'].mode()[0], inplace=True)  
/tmp/ipython-input-2612038968.py:9: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method. The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.
```

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
dataset['AgeCategory'].fillna(dataset['AgeCategory'].mode()[0], inplace=True)  
/tmp/ipython-input-2612038968.py:10: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method. The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.
```

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
dataset['Race'].fillna(dataset['Race'].mode()[0], inplace=True)  
/tmp/ipython-input-2612038968.py:11: FutureWarning: A value is trying to be set on a
```

copy of a DataFrame or Series through chained assignment using an inplace method. The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
dataset['Diabetic'].fillna(dataset['Diabetic'].mode()[0], inplace=True)
/tmp/ipython-input-2612038968.py:12: FutureWarning: A value is trying to be set on a
copy of a DataFrame or Series through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work because
the intermediate object on which we are setting values always behaves as a copy.
```

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
dataset['PhysicalActivity'].fillna(dataset['PhysicalActivity'].mode()[0], inplace=True)
/tmp/ipython-input-2612038968.py:13: FutureWarning: A value is trying to be set on a
copy of a DataFrame or Series through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work because
the intermediate object on which we are setting values always behaves as a copy.
```

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
dataset['GenHealth'].fillna(dataset['GenHealth'].mode()[0], inplace=True)
/tmp/ipython-input-2612038968.py:14: FutureWarning: A value is trying to be set on a
copy of a DataFrame or Series through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work because
the intermediate object on which we are setting values always behaves as a copy.
```

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
dataset['SleepTime'].fillna(dataset['SleepTime'].mean(), inplace=True)
/tmp/ipython-input-2612038968.py:15: FutureWarning: A value is trying to be set on a
copy of a DataFrame or Series through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work because
the intermediate object on which we are setting values always behaves as a copy.
```

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
dataset['KidneyDisease'].fillna(dataset['KidneyDisease'].mode()[0], inplace=True)
```

```
In [7]: dataset.head()
```

```
Out[7]:
```

|   | HeartDisease | BMI   | Smoking | AlcoholDrinking | Stroke | PhysicalHealth | MentalHealth | I    |
|---|--------------|-------|---------|-----------------|--------|----------------|--------------|------|
| 0 | No           | 16.60 | Yes     |                 | No     | No             | 3.0          | 30.0 |
| 1 | No           | 20.34 | No      |                 | No     | Yes            | 0.0          | 0.0  |
| 2 | No           | 26.58 | Yes     |                 | No     | No             | 20.0         | 30.0 |
| 3 | No           | 24.21 | No      |                 | No     | No             | 0.0          | 0.0  |
| 4 | No           | 23.71 | No      |                 | No     | No             | 28.0         | 0.0  |

```
In [8]: dataset.isna().sum().sum()
```

```
Out[8]: np.int64(0)
```

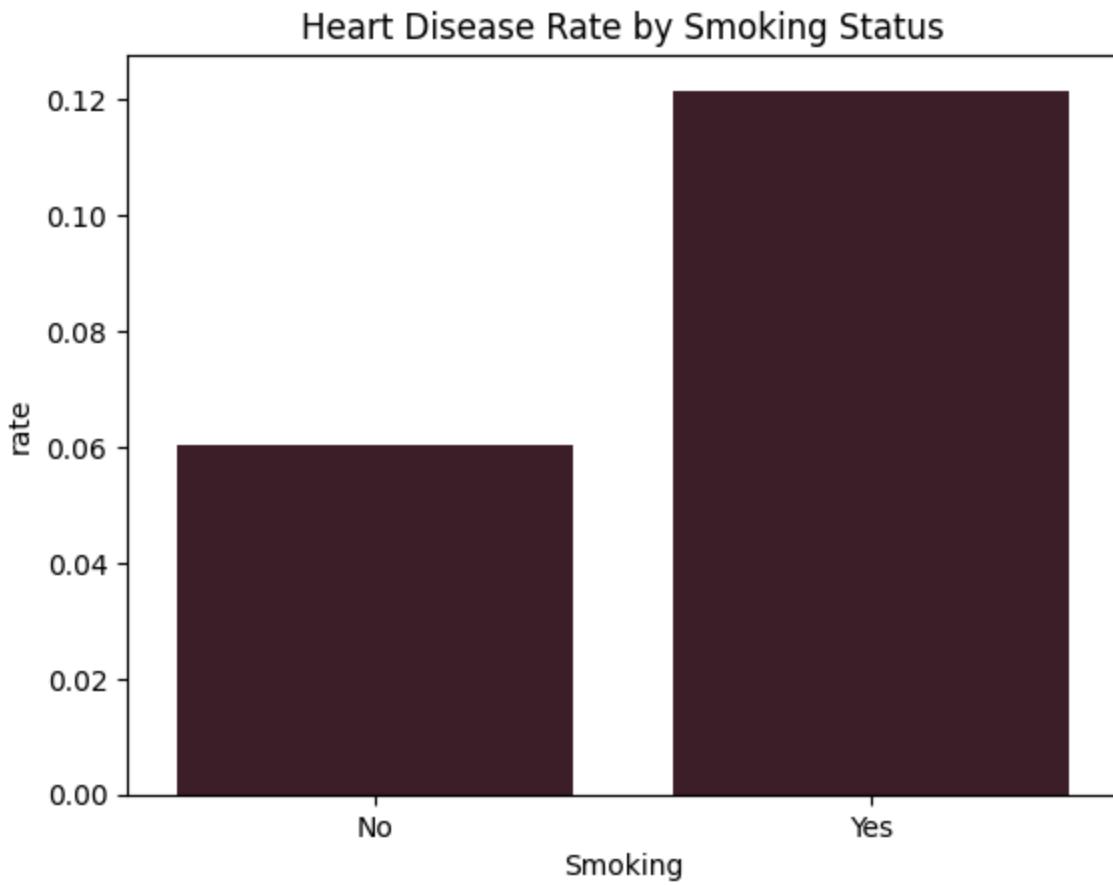
```
In [9]: dataset['HeartDisease'] = dataset['HeartDisease'].map({'Yes': 1, 'No': 0})
```

## 1. Heart disease against smoking habit

```
In [10]: heart_disease_vs_smoking_rate = (
    dataset.groupby('Smoking')['HeartDisease']
    .aggregate('mean', 'sum')
    .reset_index(name='rate')
)
print(heart_disease_vs_smoking_rate)
```

| Smoking | rate     |
|---------|----------|
| 0 No    | 0.060346 |
| 1 Yes   | 0.121574 |

```
In [11]: sns.barplot(
    data = heart_disease_vs_smoking_rate,
    x = 'Smoking',
    y = 'rate',
    color = '#451828'
)
plt.title("Heart Disease Rate by Smoking Status")
plt.show()
```



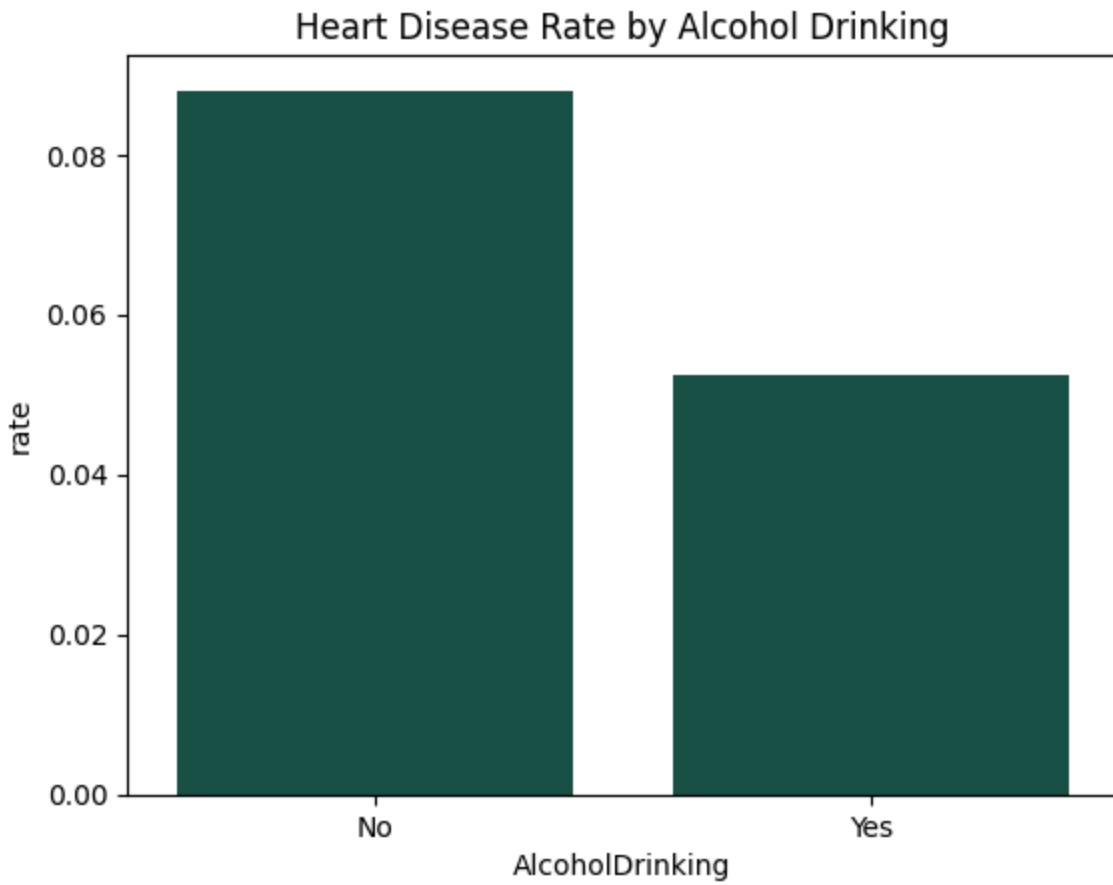
1. People who smoke have a 12.16% rate of heart disease.
2. People who don't smoke have a 6.03% rate.
3. Smoking nearly doubles the risk of heart disease in this dataset.

## 2. Heart Disease against alcohol consumption

```
In [12]: heart_disease_vs_alcohol_drinking = (
    dataset.groupby('AlcoholDrinking')[ 'HeartDisease' ]
    .aggregate('mean', 'sum')
    .reset_index(name='rate')
)
print(heart_disease_vs_alcohol_drinking)
```

| AlcoholDrinking | rate     |
|-----------------|----------|
| No              | 0.088023 |
| Yes             | 0.052361 |

```
In [13]: sns.barplot(
    data = heart_disease_vs_alcohol_drinking,
    x = 'AlcoholDrinking',
    y = 'rate',
    color = '#105e51'
)
plt.title("Heart Disease Rate by Alcohol Drinking")
plt.show()
```



1. People who do not drink alcohol have a 8.8% rate of heart disease.
2. People who do drink alcohol have a 5.24% rate.
3. Surprisingly, the dataset suggests that alcohol drinkers have a lower heart disease rate than non-drinkers.

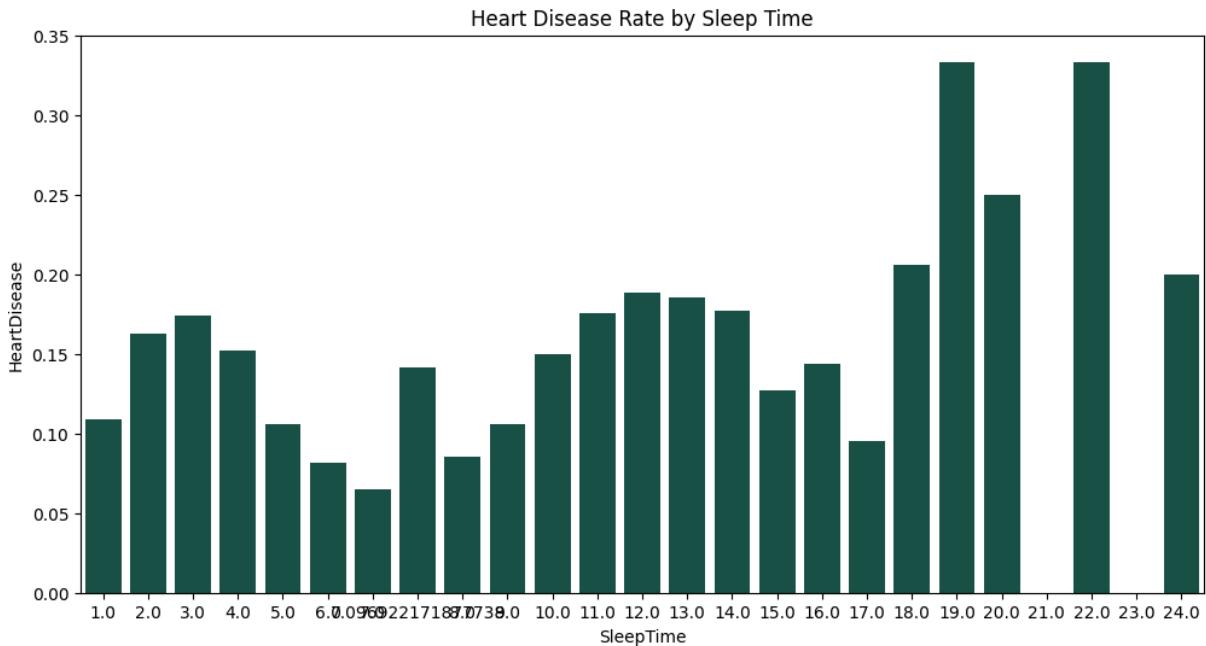
### 3. Heart Disease against sleeping pattern

```
In [14]: heart_disease_vs_sleeping_pattern = (
    dataset.groupby('SleepTime')['HeartDisease']
    .mean()
    .sort_values(ascending=False)
    .reset_index()
)

print(heart_disease_vs_sleeping_pattern)
```

|    | SleepTime | HeartDisease |
|----|-----------|--------------|
| 0  | 22.000000 | 0.333333     |
| 1  | 19.000000 | 0.333333     |
| 2  | 20.000000 | 0.250000     |
| 3  | 18.000000 | 0.205882     |
| 4  | 24.000000 | 0.200000     |
| 5  | 12.000000 | 0.188748     |
| 6  | 13.000000 | 0.185567     |
| 7  | 14.000000 | 0.176955     |
| 8  | 11.000000 | 0.175904     |
| 9  | 3.000000  | 0.174284     |
| 10 | 2.000000  | 0.162643     |
| 11 | 4.000000  | 0.152168     |
| 12 | 10.000000 | 0.149422     |
| 13 | 16.000000 | 0.144068     |
| 14 | 7.096922  | 0.141553     |
| 15 | 15.000000 | 0.126984     |
| 16 | 1.000000  | 0.108893     |
| 17 | 5.000000  | 0.105798     |
| 18 | 9.000000  | 0.105608     |
| 19 | 17.000000 | 0.095238     |
| 20 | 8.000000  | 0.085616     |
| 21 | 6.000000  | 0.081302     |
| 22 | 7.000000  | 0.064643     |
| 23 | 21.000000 | 0.000000     |
| 24 | 23.000000 | 0.000000     |

```
In [15]: plt.figure(figsize=(12,6))
sns.barplot(
    data=heart_disease_vs_sleeping_pattern,
    x='SleepTime',
    y='HeartDisease', # the mean values column
    color='#105e51'
)
plt.title("Heart Disease Rate by Sleep Time")
plt.show()
```



1. Very low or very high SleepTime (like 1–4 hours or 18–24 hours) has higher heart disease rates.
2. Middle-range sleep (6–8 hours) shows lower heart disease rates.
3. Both insufficient and excessive sleep are linked with increased heart disease risk.
4. Most common healthy sleep range (6–8 hours) correlates with lower rates, consistent with medical studies.

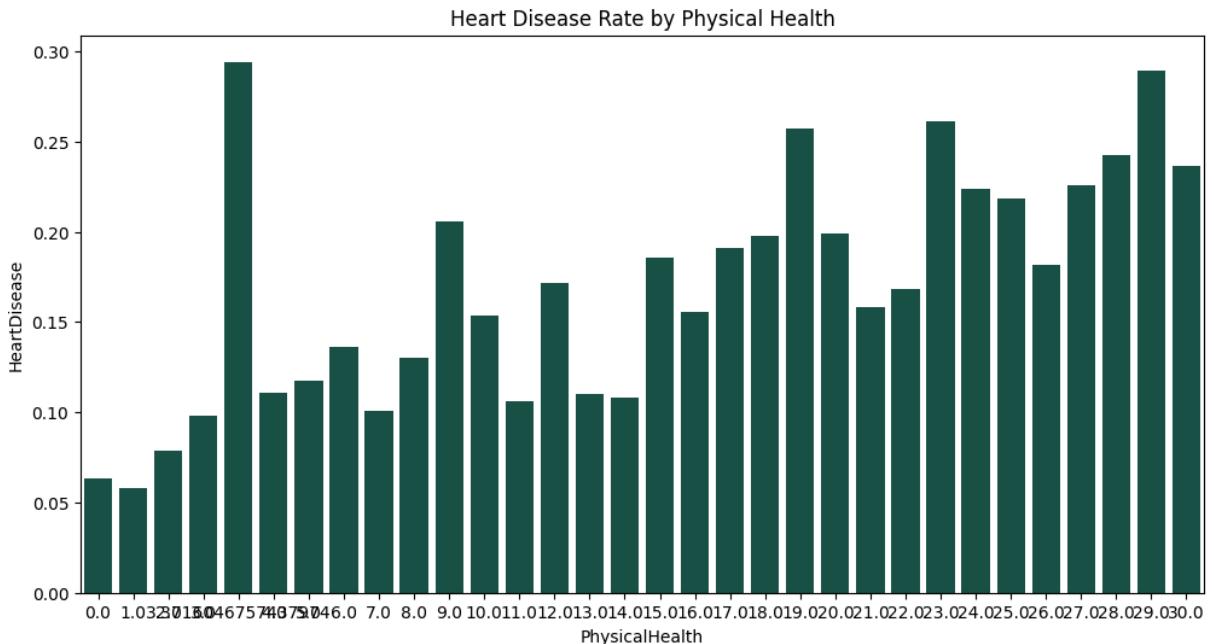
## 4. Heart Disease against Physical Health

```
In [16]: heart_disease_vs_physical_health = (
    dataset.groupby('PhysicalHealth')[ 'HeartDisease' ]
    .mean()
    .sort_values(ascending=False)
    .reset_index()
)

print(heart_disease_vs_physical_health)
```

|    | PhysicalHealth | HeartDisease |
|----|----------------|--------------|
| 0  | 3.371605       | 0.294118     |
| 1  | 29.000000      | 0.289216     |
| 2  | 23.000000      | 0.260870     |
| 3  | 19.000000      | 0.257143     |
| 4  | 28.000000      | 0.242152     |
| 5  | 30.000000      | 0.236736     |
| 6  | 27.000000      | 0.225806     |
| 7  | 24.000000      | 0.223881     |
| 8  | 25.000000      | 0.218401     |
| 9  | 9.000000       | 0.205556     |
| 10 | 20.000000      | 0.199316     |
| 11 | 18.000000      | 0.197605     |
| 12 | 17.000000      | 0.190909     |
| 13 | 15.000000      | 0.185555     |
| 14 | 26.000000      | 0.181818     |
| 15 | 12.000000      | 0.171901     |
| 16 | 22.000000      | 0.168539     |
| 17 | 21.000000      | 0.158147     |
| 18 | 16.000000      | 0.155556     |
| 19 | 10.000000      | 0.153677     |
| 20 | 6.000000       | 0.136220     |
| 21 | 8.000000       | 0.129870     |
| 22 | 5.000000       | 0.117686     |
| 23 | 4.000000       | 0.110564     |
| 24 | 13.000000      | 0.109890     |
| 25 | 14.000000      | 0.107847     |
| 26 | 11.000000      | 0.105882     |
| 27 | 7.000000       | 0.100454     |
| 28 | 3.000000       | 0.097830     |
| 29 | 2.000000       | 0.078562     |
| 30 | 0.000000       | 0.063638     |
| 31 | 1.000000       | 0.057685     |

```
In [17]: plt.figure(figsize=(12,6))
sns.barplot(
    data=heart_disease_vs_physical_health,
    x='PhysicalHealth',
    y='HeartDisease',
    color="#105e51"
)
plt.title("Heart Disease Rate by Physical Health")
plt.show()
```



1. PhysicalHealth is likely number of days in the past month with poor physical health.
2. Higher values > more unhealthy days > higher heart disease rates.
3. Low values (0–5 days) have lower risk, while 20+ days show much higher risk .

## 5. Heart Disease against various factors like Smoking, AlcoholDrinking, SleepTime, PhysicalHealth

```
In [18]: # Group by multiple features
combined_features = ['Smoking', 'AlcoholDrinking', 'SleepTime', 'PhysicalHealth']
heart_disease_combined = (
    dataset.groupby(combined_features)['HeartDisease']
    .agg(['sum', 'mean'])
    .sort_values(by='sum', ascending=False)
    .reset_index()
)

# Show top 5 combinations
top5_combined = heart_disease_combined.head(5)
print(top5_combined)
```

|   | Smoking | AlcoholDrinking | SleepTime | PhysicalHealth | sum | mean          |
|---|---------|-----------------|-----------|----------------|-----|---------------|
| 0 | Yes     |                 | No        | 8.0            | 0.0 | 2635 0.105085 |
| 1 | No      |                 | No        | 8.0            | 0.0 | 2190 0.050503 |
| 2 | Yes     |                 | No        | 7.0            | 0.0 | 1902 0.080916 |
| 3 | No      |                 | No        | 7.0            | 0.0 | 1771 0.039187 |
| 4 | Yes     |                 | No        | 6.0            | 0.0 | 1457 0.087739 |

```
In [19]: top5_combined['combo_readable'] = (
    "Smoking: " + top5_combined['Smoking'].astype(str) + ", " +
```

```

    "Alcohol: " + top5_combined['AlcoholDrinking'].astype(str) + ", " +
    "Sleep: " + top5_combined['SleepTime'].astype(str) + " hrs, " +
    "PhysicalHealth: " + top5_combined['PhysicalHealth'].astype(str) + " days"
)

# Sort by 'sum' descending to get bars in order
top5_combined = top5_combined.sort_values(by='sum', ascending=False)

```

/tmp/ipython-input-2903146847.py:1: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)  
top5\_combined['combo\_readable'] = (

In [20]: top5\_combined

|   | Smoking | AlcoholDrinking | SleepTime | PhysicalHealth | sum      | mean     | combo_readable  |
|---|---------|-----------------|-----------|----------------|----------|----------|---|
| 0 | Yes     | No              | 8.0       |                | 0.0 2635 | 0.105085 | Smoking: Yes,<br>Alcohol: No,<br>Sleep: 8.0 hrs,<br>Phys... |
| 1 | No      | No              | 8.0       |                | 0.0 2190 | 0.050503 | Smoking: No,<br>Alcohol: No,<br>Sleep: 8.0 hrs,<br>Phys...  |
| 2 | Yes     | No              | 7.0       |                | 0.0 1902 | 0.080916 | Smoking: Yes,<br>Alcohol: No,<br>Sleep: 7.0 hrs,<br>Phys... |
| 3 | No      | No              | 7.0       |                | 0.0 1771 | 0.039187 | Smoking: No,<br>Alcohol: No,<br>Sleep: 7.0 hrs,<br>Phys...  |
| 4 | Yes     | No              | 6.0       |                | 0.0 1457 | 0.087739 | Smoking: Yes,<br>Alcohol: No,<br>Sleep: 6.0 hrs,<br>Phys... |

In [21]: # Plot horizontal bar chart

```

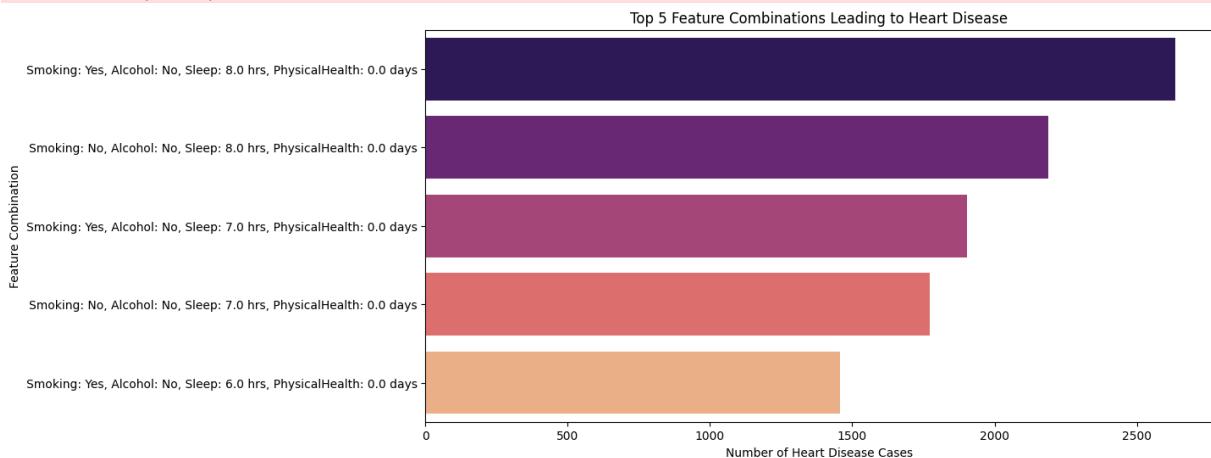
plt.figure(figsize=(12,6))
sns.barplot(
    data=top5_combined,
    x='sum',
    y='combo_readable', # readable labels
    palette='magma'
)
plt.xlabel("Number of Heart Disease Cases")
plt.ylabel("Feature Combination")

```

```
plt.title("Top 5 Feature Combinations Leading to Heart Disease")
plt.show()
```

```
/tmp/ipython-input-914638001.py:3: FutureWarning:
Passing `palette` without assigning `hue` is deprecated and will be removed in v0.1
4.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.
```

```
sns.barplot(
```



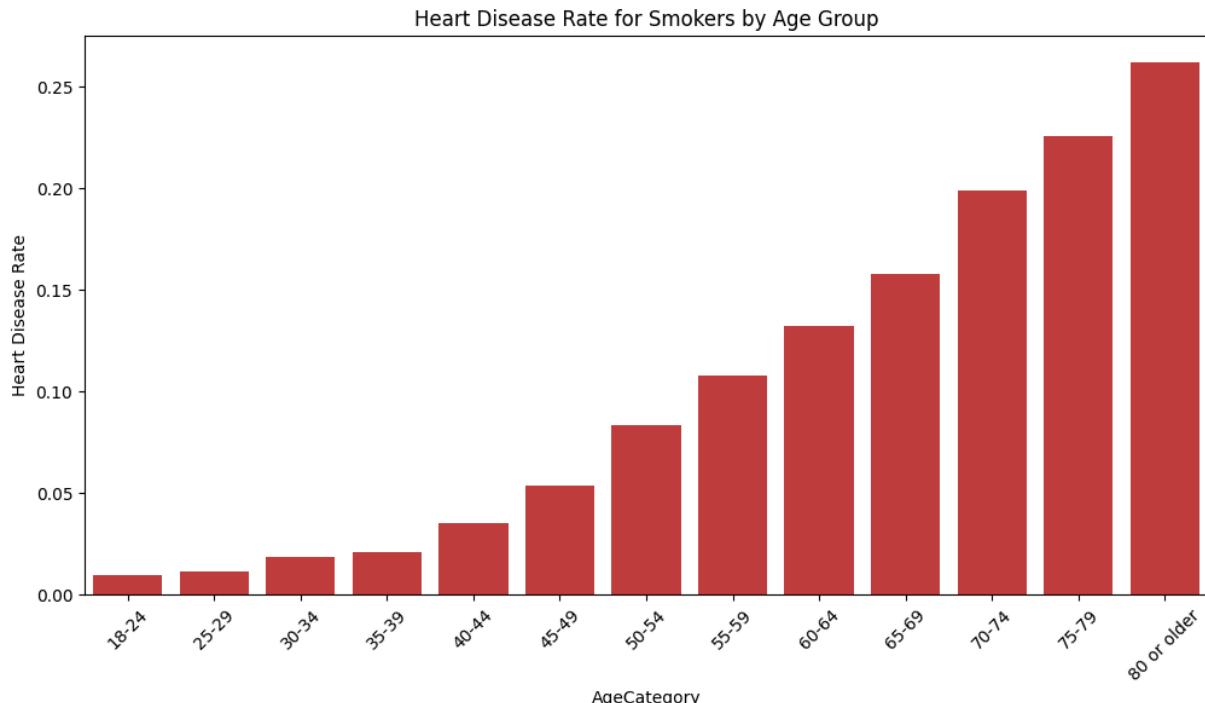
1. The highest count and risk combination is smokers who do not drink, sleep 8 hours, and have 0 poor health days: 2,633 people.
2. Second is non-smokers, non-drinkers, sleep 8 hours, 0 poor health days: 2,186 people.
3. Third is smokers, non-drinkers, sleep 7 hours, 0 poor health days, 1,901 people.
4. Smoking significantly increases risk across different sleep durations.
5. Alcohol drinking seems less prevalent in the high-risk groups here.
6. Sleep time impacts risk but not as strongly as smoking.
7. People with 0 poor physical health days still show significant risk when smoking, meaning smoking alone is a major risk factor.

## 6. Smoker Against Age

```
In [22]: smoke_age = (
    dataset[dataset['Smoking'] == 'Yes']
    .groupby('AgeCategory')['HeartDisease']
    .mean()
    .reset_index()
    .sort_values('AgeCategory')
)
print(smoke_age)
```

|    | AgeCategory | HeartDisease |
|----|-------------|--------------|
| 0  | 18-24       | 0.009747     |
| 1  | 25-29       | 0.011454     |
| 2  | 30-34       | 0.018555     |
| 3  | 35-39       | 0.021067     |
| 4  | 40-44       | 0.035128     |
| 5  | 45-49       | 0.053307     |
| 6  | 50-54       | 0.083475     |
| 7  | 55-59       | 0.107884     |
| 8  | 60-64       | 0.132102     |
| 9  | 65-69       | 0.157847     |
| 10 | 70-74       | 0.198914     |
| 11 | 75-79       | 0.225551     |
| 12 | 80 or older | 0.261882     |

```
In [23]: plt.figure(figsize=(12,6))
sns.barplot(
    data=smoke_age,
    x='AgeCategory',
    y='HeartDisease',
    color='#d62728'
)
plt.title("Heart Disease Rate for Smokers by Age Group")
plt.ylabel("Heart Disease Rate")
plt.xticks(rotation=45)
plt.show()
```



## 7. Drinker Against Age

```
In [24]: drink_age = (
    dataset[dataset['AlcoholDrinking'] == 'Yes']
    .groupby('AgeCategory')['HeartDisease']
    .mean()
```

```

    .reset_index()
    .sort_values('AgeCategory')
)
print(drink_age)

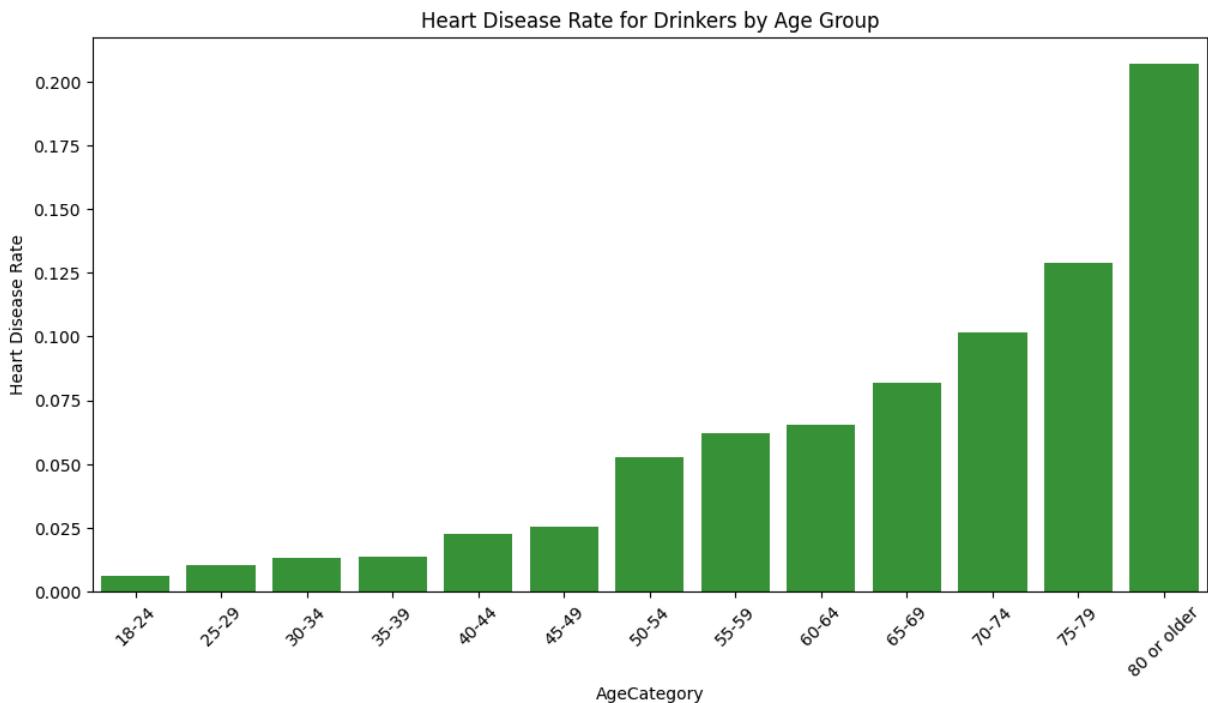
```

|    | AgeCategory | HeartDisease |
|----|-------------|--------------|
| 0  | 18-24       | 0.005917     |
| 1  | 25-29       | 0.010185     |
| 2  | 30-34       | 0.013308     |
| 3  | 35-39       | 0.013699     |
| 4  | 40-44       | 0.022715     |
| 5  | 45-49       | 0.025550     |
| 6  | 50-54       | 0.052578     |
| 7  | 55-59       | 0.062218     |
| 8  | 60-64       | 0.065553     |
| 9  | 65-69       | 0.081905     |
| 10 | 70-74       | 0.101440     |
| 11 | 75-79       | 0.128889     |
| 12 | 80 or older | 0.207048     |

```

In [25]: plt.figure(figsize=(12,6))
sns.barplot(
    data=drink_age,
    x='AgeCategory',
    y='HeartDisease',
    color="#2ca02c"
)
plt.title("Heart Disease Rate for Drinkers by Age Group")
plt.ylabel("Heart Disease Rate")
plt.xticks(rotation=45)
plt.show()

```



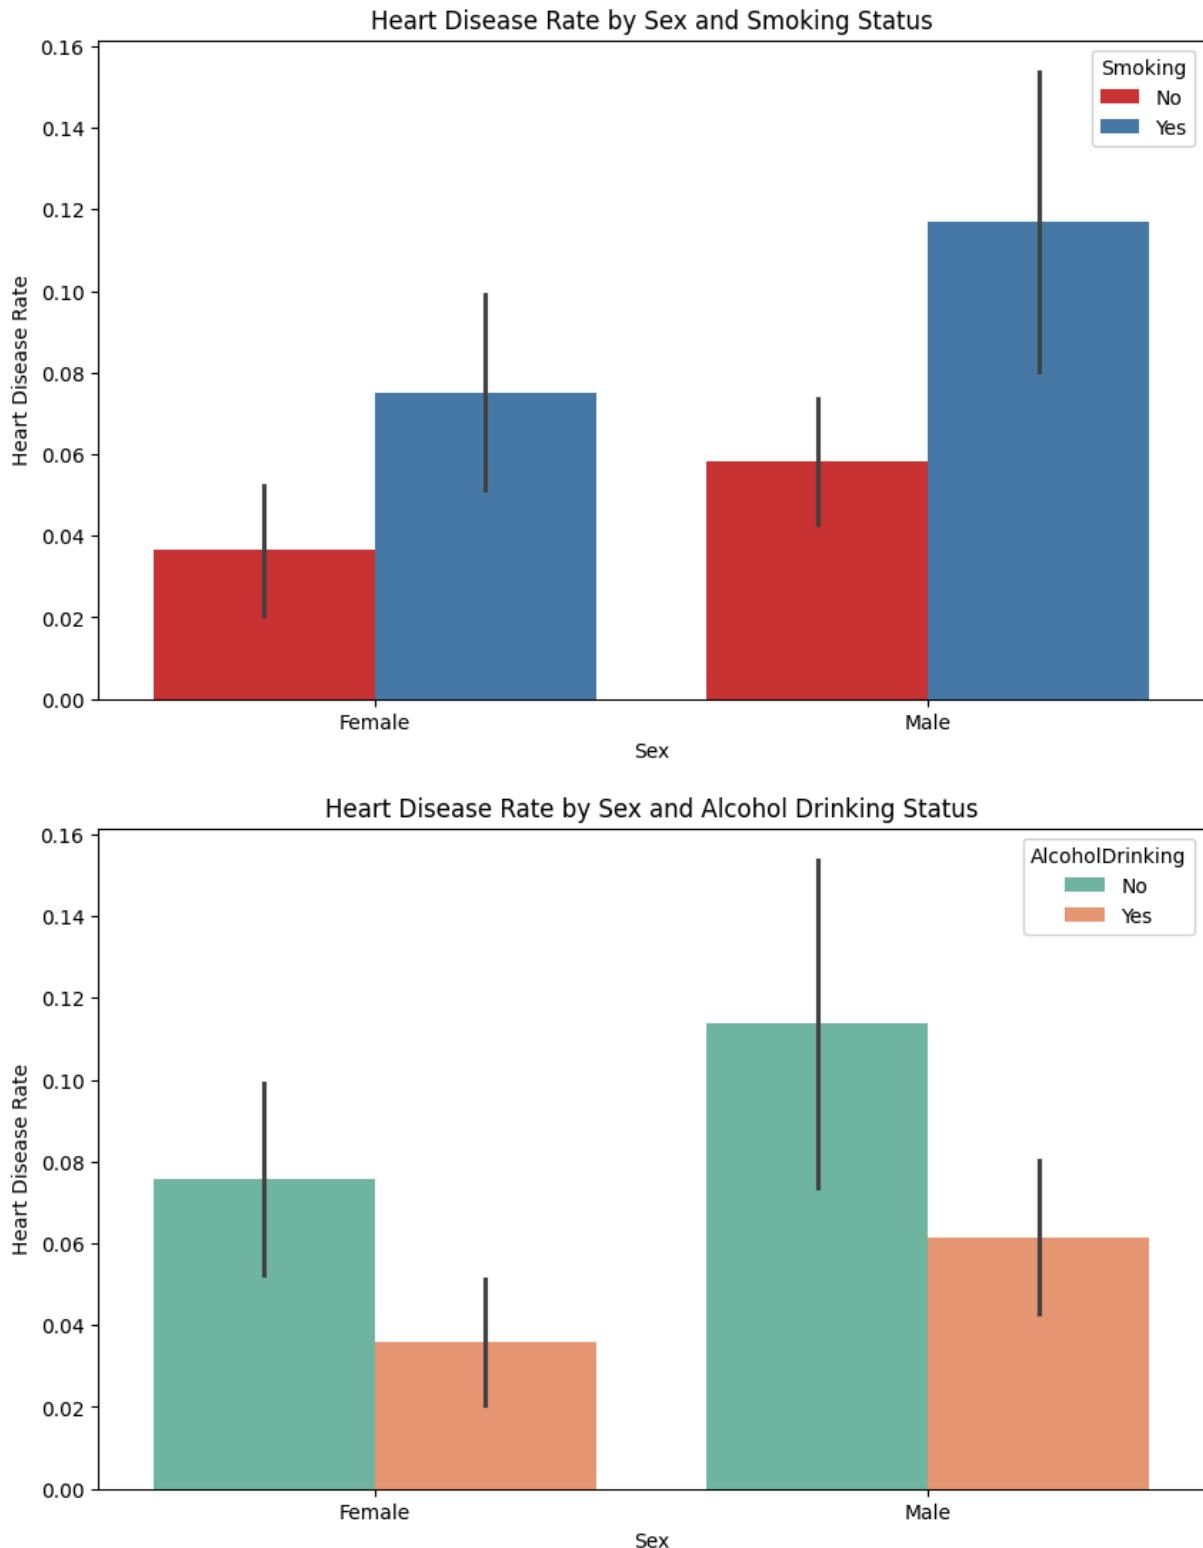
## 8. Smoker and Drinker Against Sex

```
In [26]: combo = (
    dataset
    .groupby(['Smoking', 'AlcoholDrinking', 'Sex'])['HeartDisease']
    .mean()
    .reset_index()
)
print(combo)
```

|   | Smoking | AlcoholDrinking | Sex    | HeartDisease |
|---|---------|-----------------|--------|--------------|
| 0 | No      | No              | Female | 0.052399     |
| 1 | No      | No              | Male   | 0.073667     |
| 2 | No      | Yes             | Female | 0.020619     |
| 3 | No      | Yes             | Male   | 0.042723     |
| 4 | Yes     | No              | Female | 0.099058     |
| 5 | Yes     | No              | Male   | 0.153744     |
| 6 | Yes     | Yes             | Female | 0.051194     |
| 7 | Yes     | Yes             | Male   | 0.080240     |

```
In [27]: plt.figure(figsize=(10,6))
sns.barplot(
    data=combo,
    x='Sex',
    y='HeartDisease',
    hue='Smoking',
    palette='Set1'
)
plt.title("Heart Disease Rate by Sex and Smoking Status")
plt.ylabel("Heart Disease Rate")
plt.show()

plt.figure(figsize=(10,6))
sns.barplot(
    data=combo,
    x='Sex',
    y='HeartDisease',
    hue='AlcoholDrinking',
    palette='Set2'
)
plt.title("Heart Disease Rate by Sex and Alcohol Drinking Status")
plt.ylabel("Heart Disease Rate")
plt.show()
```



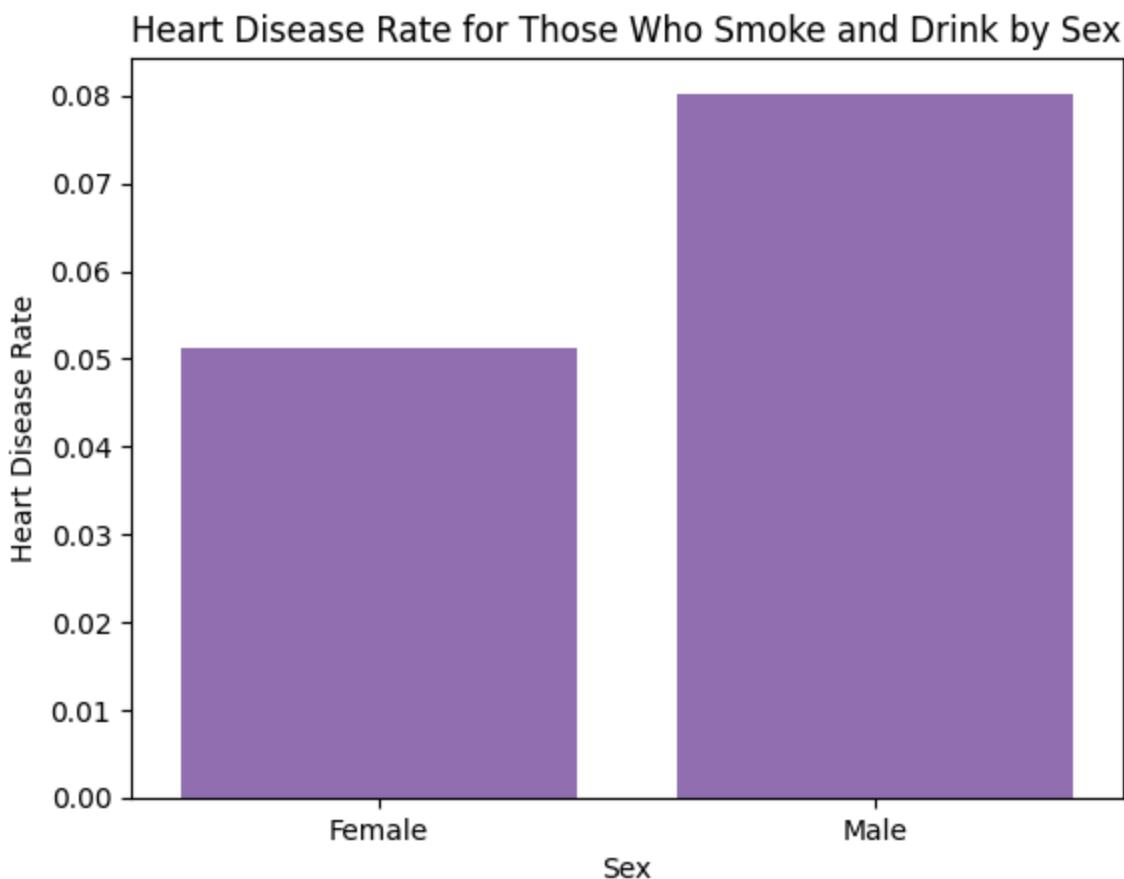
## 9. Combined smoker+drinker vs. sex

```
In [28]: both_smoke_drink = (
    dataset[(dataset['Smoking'] == 'Yes') & (dataset['AlcoholDrinking'] == 'Yes')]
    .groupby('Sex')['HeartDisease']
    .mean()
    .reset_index()
```

```
)  
print(both_smoke_drink)
```

```
Sex  HeartDisease  
0  Female      0.051194  
1  Male        0.080240
```

```
In [29]: sns.barplot(  
    data=both_smoke_drink,  
    x='Sex',  
    y='HeartDisease',  
    color='#9467bd'  
)  
plt.title("Heart Disease Rate for Those Who Smoke and Drink by Sex")  
plt.ylabel("Heart Disease Rate")  
plt.show()
```



This study analyzed the impact of lifestyle and health factors on heart disease using data analysis and visualization techniques. Smoking was found to be the strongest risk factor, with smokers showing nearly double the heart disease rate compared to non-smokers. Sleep duration showed a U-shaped relationship, where both insufficient and excessive sleep increased risk, while 6-8 hours was associated with lower rates. Poor physical health days were strongly correlated with higher heart disease risk. Combined feature analysis confirmed that smoking consistently appears in the highest-risk groups. Overall, the results demonstrate how data-driven decision support systems can identify key risk patterns for cardiovascular disease.

# Random Forest Classifier

```
In [30]: from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
from sklearn.ensemble import RandomForestClassifier
from sklearn.preprocessing import LabelEncoder
from imblearn.over_sampling import SMOTE
import pandas as pd
import numpy as np
```

```
In [31]: # 1. DATA PREPARATION - PROPER ENCODING
df = dataset.copy()

# Handle missing values for numerical columns
numerical_cols = df.select_dtypes(include=['int64', 'float64']).columns
df[numerical_cols] = df[numerical_cols].fillna(df[numerical_cols].mean())

# Handle missing values for categorical columns
categorical_cols = df.select_dtypes(include=['object', 'category']).columns
for col in categorical_cols:
    df[col].fillna(df[col].mode()[0], inplace=True)
```

/tmp/ipython-input-339661290.py:11: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method. The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df[col].fillna(df[col].mode()[0], inplace=True)
```

```
In [32]: # Encode ALL categorical variables BEFORE splitting
le = LabelEncoder()
for col in categorical_cols:
    df[col] = le.fit_transform(df[col].astype(str))

print(f"Original dataset shape: {df.shape}")
print(f"Columns after encoding: {df.columns.tolist()}\n")
```

Original dataset shape: (319795, 18)  
Columns after encoding: ['HeartDisease', 'BMI', 'Smoking', 'AlcoholDrinking', 'Stroke', 'PhysicalHealth', 'MentalHealth', 'DiffWalking', 'Sex', 'AgeCategory', 'Race', 'Diabetic', 'PhysicalActivity', 'GenHealth', 'SleepTime', 'Asthma', 'KidneyDisease', 'SkinCancer']

```
In [33]: # Verify no object types remain
print(f"Data types after encoding:\n{df.dtypes.value_counts()}\n")
```

```
Data types after encoding:
```

```
int64      14
float64     4
Name: count, dtype: int64
```

```
In [34]: # 2. PREPARE X and y
X = df.drop('HeartDisease', axis=1)
y = df['HeartDisease']

print(f"Features shape: {X.shape}")
print(f"Class distribution:\n{y.value_counts()}")
print(f"Imbalance ratio: {y.value_counts()[0] / y.value_counts()[1]:.2f}\n")
```

```
Features shape: (319795, 17)
```

```
Class distribution:
```

```
HeartDisease
```

```
0    292422
1    27373
```

```
Name: count, dtype: int64
```

```
Imbalance ratio: 10.68:1
```

```
In [35]: # 3. TRAIN-TEST SPLIT
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)

print(f"Training set size: {X_train.shape}")
print(f"Test set size: {X_test.shape}\n")
```

```
Training set size: (255836, 17)
```

```
Test set size: (63959, 17)
```

```
In [40]: # 4. APPLY SMOTE TO BALANCE CLASSES
smote = SMOTE(sampling_strategy=1.0, random_state=42, k_neighbors=3)
X_train_res, y_train_res = smote.fit_resample(X_train, y_train)

print(f"After SMOTE:")
print(f"  Class 0 (No Disease): {(y_train_res == 0).sum()}")
print(f"  Class 1 (Heart Disease): {(y_train_res == 1).sum()}")
print(f"  Total: {len(y_train_res)}\n")
```

```
After SMOTE:
```

```
  Class 0 (No Disease): 233938
  Class 1 (Heart Disease): 233938
  Total: 467876
```

```
In [41]: # 5. RANDOM FOREST WITH OPTIMIZED PARAMETERS
rf_model = RandomForestClassifier(
    n_estimators=500,
    max_depth=20,
    min_samples_split=10,
    min_samples_leaf=5,
    max_features='sqrt',
```

```

        class_weight='balanced',
        random_state=42,
        n_jobs=-1,
        bootstrap=True,
        max_samples=0.85
    )

print("Training Random Forest with SMOTE-balanced data...")
# rf_model.fit(X_train_res, y_train_res)
rf_model.fit(X_train, y_train)
print("Training complete!\n")

```

Training Random Forest with SMOTE-balanced data...  
Training complete!

In [42]: # 6. PREDICTIONS  
y\_pred\_proba = rf\_model.predict\_proba(X\_test)[:, 1]  
y\_pred = rf\_model.predict(X\_test)

In [44]: # 7. EVALUATION - DEFAULT THRESHOLD  
print("=" \* 70)  
print("RESULTS AT DEFAULT THRESHOLD (0.5)")  
print("=" \* 70)  
print(f"\nAccuracy: {accuracy\_score(y\_test, y\_pred):.4f}")  
print("\nConfusion Matrix:")  
cm = confusion\_matrix(y\_test, y\_pred)  
print(cm)  
print("\nClassification Report:")  
print(classification\_report(y\_test, y\_pred, target\_names=['No Disease', 'Heart Disease']))

=====  
RESULTS AT DEFAULT THRESHOLD (0.5)  
=====

Accuracy: 0.8165

Confusion Matrix:  
[[48779 9705]  
 [ 2032 3443]]

|               | precision | recall | f1-score | support |
|---------------|-----------|--------|----------|---------|
| No Disease    | 0.96      | 0.83   | 0.89     | 58484   |
| Heart Disease | 0.26      | 0.63   | 0.37     | 5475    |
| accuracy      |           |        | 0.82     | 63959   |
| macro avg     | 0.61      | 0.73   | 0.63     | 63959   |
| weighted avg  | 0.90      | 0.82   | 0.85     | 63959   |

In [ ]: !jupyter nbconvert --to webpdf "D:/Github/Data-Science-And-Machine-Learning-Course/

In [ ]: